

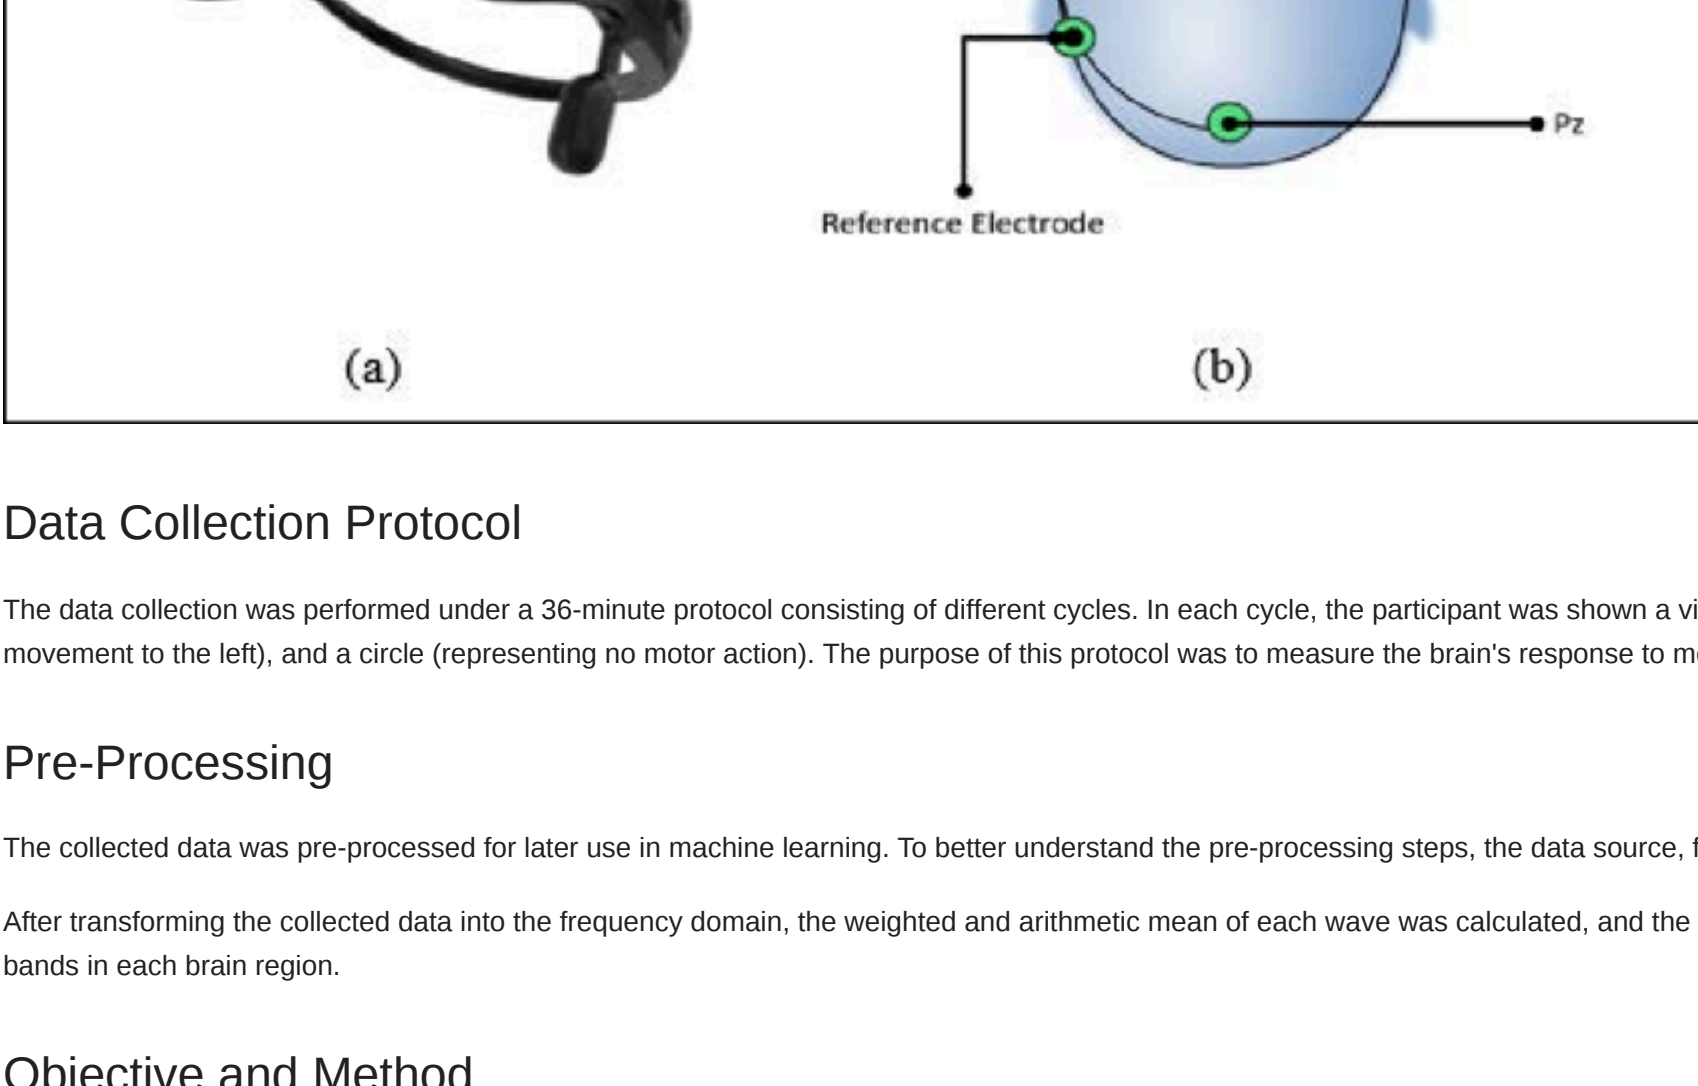
Introduction

In many parts of the world, individuals become paralyzed due to various reasons, particularly spinal cord injuries, leaving them unable to move their limbs. Machine learning models developed using EEG datasets have the potential to help these patients communicate with the outside world through Brain-Computer Interfaces (BCI) and even control mechanical limbs. This project focuses on analyzing brain waves using EEG (Electroencephalography) data and performing classification based on these signals. EEG is a technique that measures the brain's electrical activity and is commonly employed in the analysis of neurological and psychological conditions. In this study, I utilized data obtained from the EMOVTv Insight + 5 Channel EEG device to examine the frequency bands of different brain regions, specifically focusing on power values associated with components such as alpha, beta, delta, theta, and gamma waves.

In this project, I tested the classification of different hand movements using EEG data and achieved **98% accuracy** using the **KNN model** on the test dataset. As more advanced models are developed and more complex movements are classified, the possibility of creating a fully controlled BCI becomes increasingly achievable.

Dataset

The dataset consists of data obtained from the EMOVTv Insight + 5 Channel EEG device. This device has a sampling frequency of 128 Hz and uses a 16-bit analog-to-digital converter to measure brain waves. The EEG device records signals from the following brain regions: AF3, T7, T8, Pz. The dataset includes raw EEG data associated with visual stimuli representing different motor actions. During the data collection process, visual stimuli were shown to the participant, representing motor actions. While the EEG device recorded the brain waves related to these stimuli. In the situation where the participant's eyes are closed, they were given an auditory signal at specific times to open their eyes and view the images.



Data Collection Protocol

The data collection was performed under a 36-minute protocol consisting of different cycles. In each cycle, the participant was shown a visual stimulus: an arrow to the right (representing a movement to the right), an arrow to the left (representing a movement to the left), and a circle (representing no motor action). The purpose of this protocol was to measure the brain's response to motor movements and to make this data generalizable for machine learning models.

Pre-Processing

The collected data was pre-processed for later use in machine learning. To better understand the pre-processing steps, the data source, format, and characteristics are explained below.

After transforming the collected data into the frequency domain, the weighted and arithmetic mean of each wave was calculated, and the resulting matrix was created with the dimensions of 5x5. This matrix represents the power of different frequency bands in each brain region.

Objective and Method

The goal of this project is to classify brain activities using EEG data. For this purpose, classification will be performed based on the power of the frequency bands from each brain region, and the best model will be selected by evaluating the accuracy of this classification. The following steps will be taken:

Data Pre-processing: The raw EEG data will be transformed into the frequency domain, followed by normalization and cleaning.

Modeling: Different classifiers (e.g., KNN, Random Forest, Logistic Regression, etc.) will be used to create models, and hyperparameter optimization will be conducted.

Model Evaluation: Model performance will be evaluated using accuracy and other metrics.

Goals

This project aims to classify brain activities through the analysis of EEG data. To achieve this goal, various classifiers will be applied, and the accuracy of each model will be compared. Additionally, the pre-processing phase and modeling process will be carried out to obtain the best possible results.

This machine learning model was developed using the dataset titled "Brain wave data from hands movement of EEG" available on the Kaggle platform.

The model was created by **CaMaya Elk**.

You can access the dataset here:

Brain wave data from hands movement of EEG

1. Importing Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import os

from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, KFold, cross_val_score, GridSearchCV
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.ensemble import AdaBoostClassifier, GradientBoostingClassifier, RandomForestClassifier

import warnings
warnings.filterwarnings('ignore')
```

2. Merging EEG Data from Five Distinct Participants

```
csv_files = ["P001.csv", "P002.csv", "P003.csv", "P004.csv", "P005.csv"]
dataframes = []

for file in csv_files:
    df = pd.read_csv(file)
    dataframes.append(df)

merged_df = pd.concat(dataframes, ignore_index=True)
merged_df.to_csv("merged_data.csv", index=False)

print("All CSV files were merged row by row and saved as 'merged_data.csv'")
All CSV files were merged row by row and saved as 'merged_data.csv'!
```

3. Describing the Data

```
df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 548, dtype: int64
Data columns (total 26 columns):
 #   Values  POW.AF3.Theta  POW.AF3.Alpha  POW.AF3.Beta  POW.AF3.BetaH  POW.AF3.Gamma  POW.T7.Theta  POW.T7.Alpha  POW.T7.Beta  POW.T7.BetaH  ...  POW.T8.Theta  POW.T8.Alpha  POW.T8.Beta  POW.T8.BetaH  POW.T8.Gamma  POW.AF4.Theta
0    0.0      15.037535      3.466751      5.527066      2.220269      0.449631      5.655759      1.696389      0.747850      ...      3.661616      1.550242      3.410099      1.674853      0.837529      6.559598
1    0.0      10.170770      3.000238      4.004105      1.823901      0.474929      5.518808      1.997300      0.638174      ...      0.936994      1.826468      1.672080      0.848875      7.169808
2    0.0      18.103280      4.140787      3.902328      1.611893      0.511822      2.435185      0.778720      0.564740      ...      4.177790      2.200124      1.906458      1.570778      0.837001      7.692561
3    0.0      13.740777      4.113275      3.000238      1.342602      0.567074      4.280084      3.054517      0.729953      ...      2.432153      ...      2.432157      2.432157      0.800211      8.301702
...    ...    ...    ...    ...    ...    ...    ...    ...    ...    ...    ...    ...    ...    ...    ...    ...
9    0.0      1.408678      2.249824      2.249824      1.881372      0.426211      3.401131      0.742631      0.877360      ...      0.403724      2.397638      2.027718      1.232775      0.790205      6.868889
10   0.0      15.950205      3.332325      2.165254      1.152628      0.576603      2.231369      1.246149      0.496951      ...      2.081000      1.232466      1.292420      1.304634      0.782051      8.761477
11   0.0      15.493847      2.767034      2.264643      1.379978      0.616193      3.133281      4.673981      1.614407      ...      0.453968      ...      3.283734      2.978233      2.148399      1.083774      8.105113      8.222730
12   0.0      10.827897      2.344624      2.697231      1.402094      0.785339      3.372255      4.603820      2.099072      ...      0.547973      ...      3.693718      2.548129      2.700610      1.216743      0.802396      7.410793
13   0.0      7.762703      2.045052      2.997248      1.607601      0.426304      4.127768      2.595423      0.702211      ...      0.466119      1.973873      3.463463      1.527780      0.881318      6.349564
14   0.0      5.398981      1.918086      3.402179      1.809564      1.057900      4.916137      3.555354      2.946026      ...      0.933536      ...      3.467055      1.415493      4.215637      1.984765      1.055077      5.779400
15 rows x 26 columns
```

```
df.describe()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 548, dtype: int64
Data columns (total 26 columns):
 #   Values  POW.AF3.Theta  POW.AF3.Alpha  POW.AF3.Beta  POW.AF3.BetaH  POW.AF3.Gamma  POW.T7.Theta  POW.T7.Alpha  POW.T7.Beta  POW.T7.BetaH  ...  POW.T8.Theta  POW.T8.Alpha  POW.T8.Beta  POW.T8.BetaH  POW.T8.Gamma  POW.AF4.Theta
count  548.000000  548.000000  548.000000  548.000000  548.000000  548.000000  548.000000  548.000000  548.000000  548.000000  ...  548.000000  548.000000  548.000000  548.000000  548.000000  548.000000
mean    0.00405  26.742049  3.000238  1.611893  0.511822  2.435185  0.747850  3.054517  0.564740  0.496951  ...  2.081829  1.414327  1.454672  1.674853  7.613825  5.716433
std     0.897458  282.403085  61.622018  105.785738  82.960331  78.303374  324.539632  171.184032  99.810551  80.147890  ...  292.637743  78.737899  41.690065  24.528389  17.777310  29.4
min     0.000000  0.000000  0.000000  0.000000  0.000000  0.000000  0.000000  0.000000  0.000000  0.000000  ...  0.000000  0.000000  0.000000  0.000000  0.000000  0.000000
25%    0.000000  3.208777  2.166011  1.113366  0.509765  0.376574  0.695786  0.577096  0.459462  0.401296  ...  0.150972  1.738433  1.054476  0.810944  0.370133  7.
50%    0.000000  7.991018  4.407033  1.725258  1.231472  0.548427  1.427271  0.988115  0.789446  0.756512  ...  0.303749  2.956051  1.767813  1.363870  0.678196  8.
75%    2.000000  71.740446  26.737399  14.79331  2.599694  1.266578  8.697760  4.287541  2.813323  4.140122  ...  5.742385  7.892247  4.130129  3.738178  2.619889  6.
max     2.000000  1766.017548  6095.287109  2427.682617  1587.170571  1085.472295  13075.584961  7861.771484  2688.726563  1809.549316  ...  25345.798575  5625.639371  3863.484131  997.893860  850.612610  16524
```

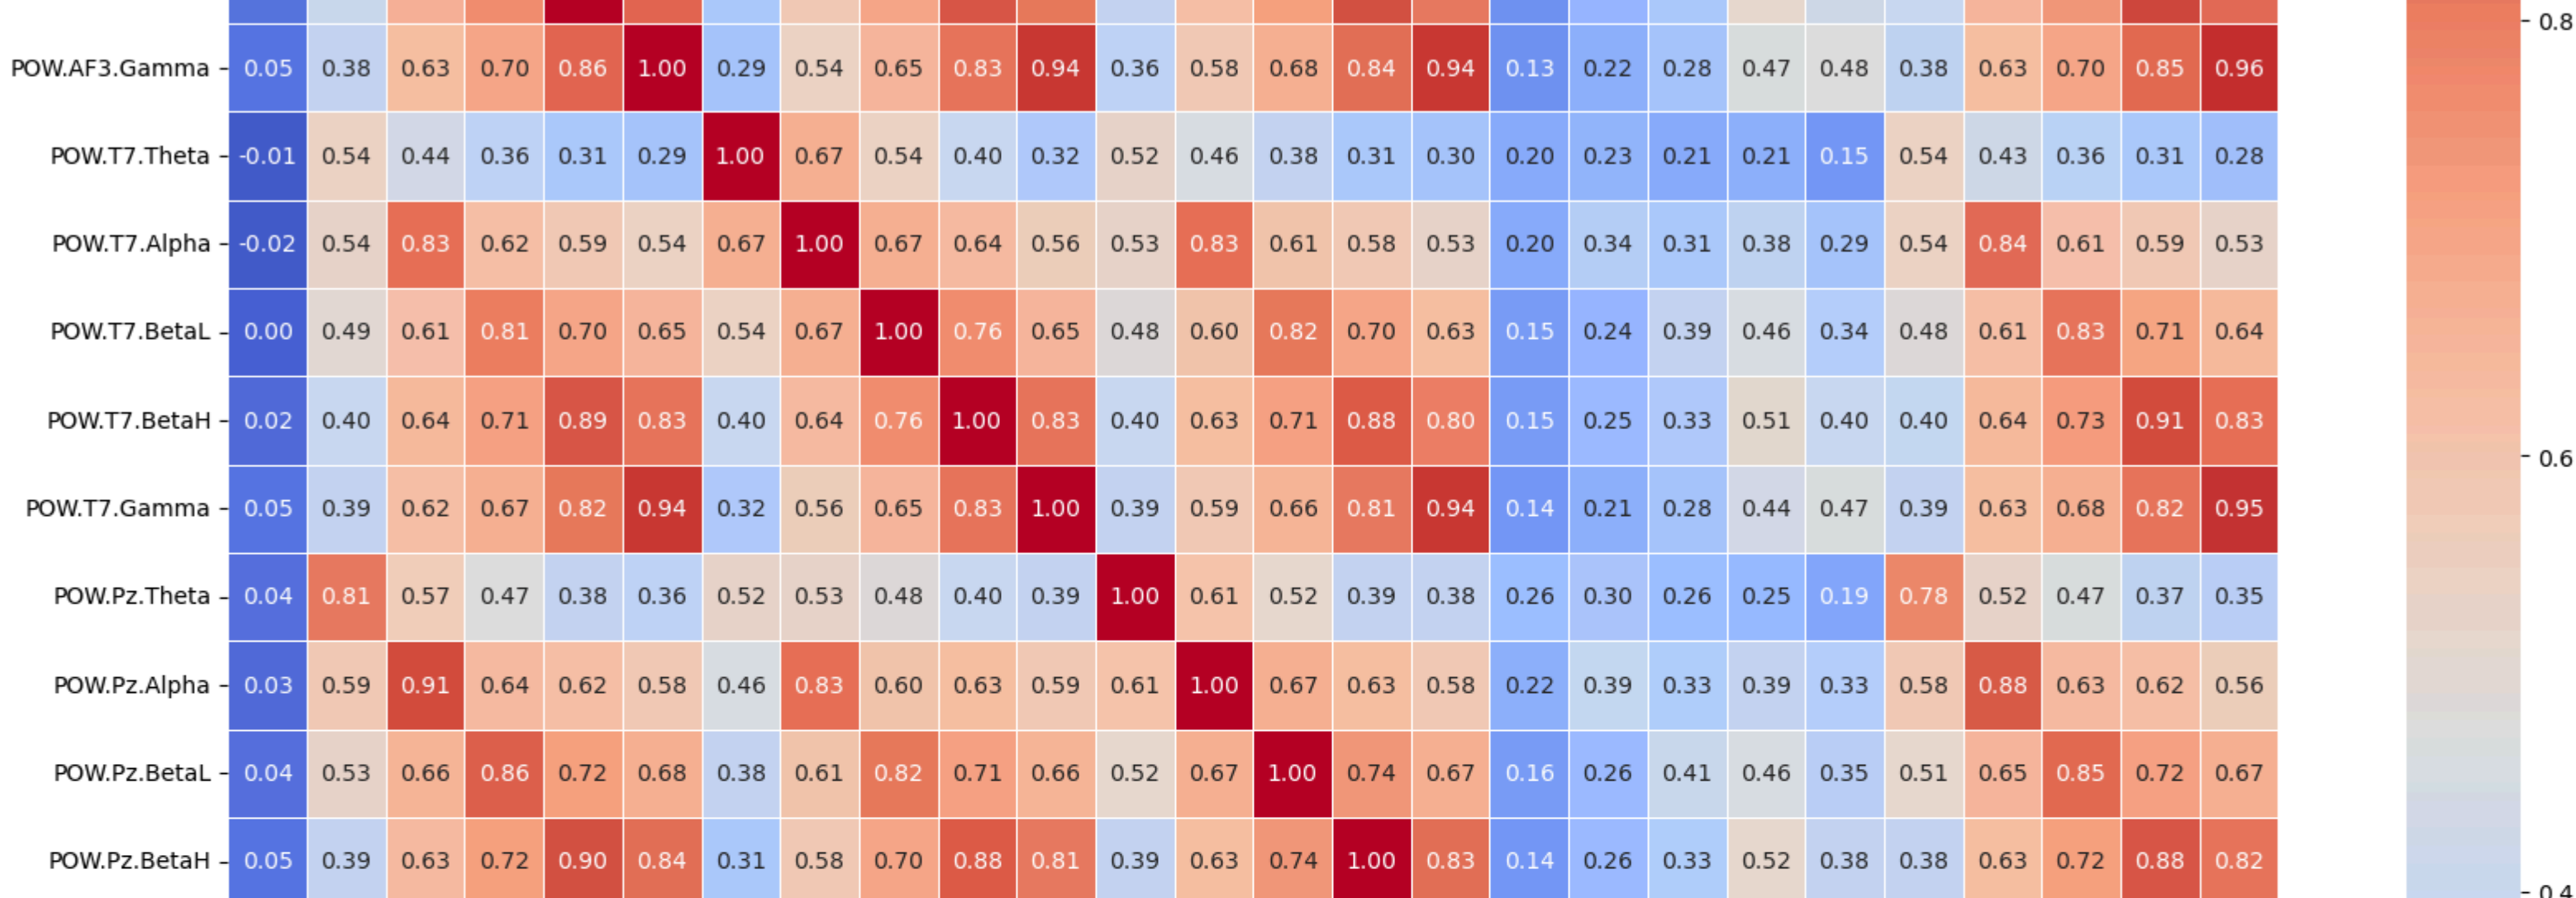
8 rows x 26 columns

The data shows significant variability, with high standard deviations indicating dynamic brain activity. Theta and Alpha bands generally have higher mean values, suggesting involvement in cognitive and attentional processes. The wide range in power values, especially in Beta and Gamma bands, may indicate different mental states or external influences. Further analysis and visualization could help interpret these patterns more effectively.

4. Plotting Correlation Heatmap

```
def plot_correlation_heatmap(df):
    corr_matrix = df.corr()
    plt.figure(figsize=(10, 10))
    sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt='.2f', linewidths=0.5)
    plt.title("Correlation of Features")
    plt.savefig("correlation_features.png")
    plt.show()

plot_correlation_heatmap(df)
```



There are some high correlations seen between different brain regions, such as AF3 Gamma and AF3 Alpha, or T7 Alpha and AF4 Alpha.

5. Outlier Detection (IQR Method)

```
def detect_outliers_iqr(df):
    outlier_indices = []
    for col in df.columns:
        Q1 = df[col].quantile(0.25)
        Q3 = df[col].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
        outlier_indices.extend(df[(df[col] < lower_bound) | (df[col] > upper_bound)].index)
    outlier_indices.extend(outlier_indices)
    df.drop(outlier_indices, inplace=True)
    return df

outliers_df = df[outlier_indices]
df_cleaned = df.drop(outlier_indices, inplace=True)

df_cleaned.to_csv("cleaned_data.csv", index=False)
```

Outliers were detected and removed from the dataset using the Interquartile Range (IQR) method. The original dataset, which had 2,222.61 entries, was cleaned, reducing it to 1,454.102 entries after removing the outliers.

Separating features and target variables (train/test) / Standardization

```
X = df_cleaned.drop(["Values"], axis=1)
y = df_cleaned["Values"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)

 scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

6. Defining Models for Evaluation

```
def get_base_models():
    models = [
        ("LR", LogisticRegression()),
        ("DT", DecisionTreeClassifier()),
        ("KNN", KNeighborsClassifier()),
        ("AdaB", AdaBoostClassifier()),
        ("GBM", GradientBoostingClassifier()),
        ("RF", RandomForestClassifier(n_estimators=50)),
        ("SVM", SVC(kernel='linear'))
    ]
    return base_models
```

7. Model Evaluation (Cross-validation)

```
def evaluate_models(X_train, y_train, models):
    names = []
    for name, model in models:
        kf = KFold(n_splits=5, shuffle=True, random_state=42)
        cv_results = cross_val_score(model, X_train, y_train, cv=kf, scoring='accuracy')
        results.append(cv_results)
        names.append(name)
        print(f'{name}: {cv_results.mean():.3f} std: {cv_results.std():.3f}')
    return names, results

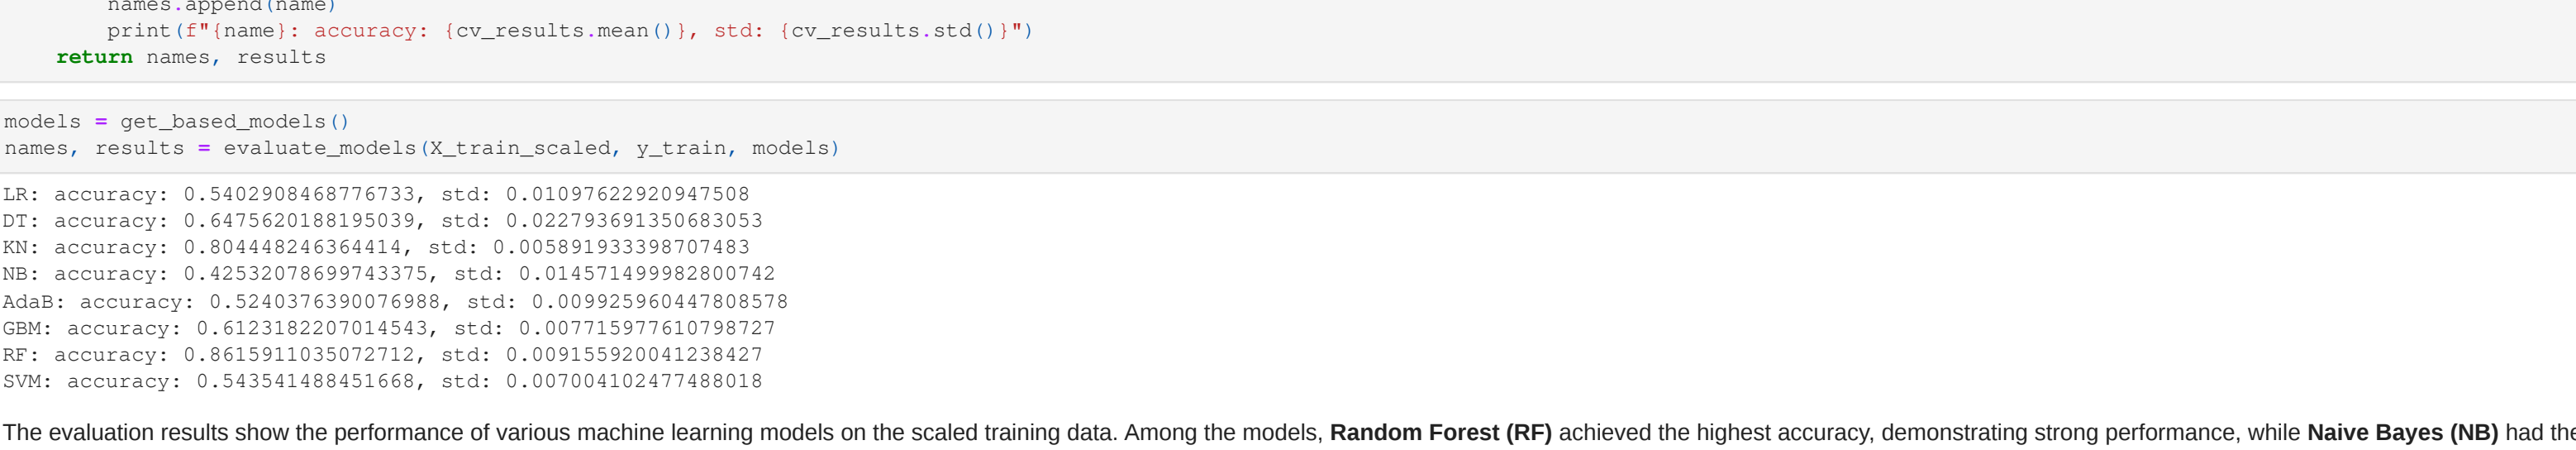
names, results = evaluate_models(X_train_scaled, y_train, models)
```

```
LR: accuracy: 0.540290488776713, std: 0.0119742202907508
DT: accuracy: 0.647620188130039, std: 0.0207916813008033
KNN: accuracy: 0.80448024634414, std: 0.0095819333970783
AdaB: accuracy: 0.420250769157575, std: 0.0104714598820742
GBM: accuracy: 0.520276376007897, std: 0.009252560467808578
RF: accuracy: 0.912321627017157, std: 0.0071514598820742
SVM: accuracy: 0.561591103077319, std: 0.00913592041238423
LR: accuracy: 0.540290488776713, std: 0.0119742202907508
DT: accuracy: 0.647620188130039, std: 0.0207916813008033
KNN: accuracy: 0.80448024634414, std: 0.0095819333970783
AdaB: accuracy: 0.420250769157575, std: 0.0104714598820742
GBM: accuracy: 0.520276376007897, std: 0.009252560467808578
RF: accuracy: 0.912321627017157, std: 0.0071514598820742
SVM: accuracy: 0.561591103077319, std: 0.00913592041238423
```

The evaluation results show the performance of various machine learning models on the scaled training data. Among the models, **Random Forest (RF)** achieved the highest accuracy, demonstrating strong performance, while **Naive Bayes (NB)** had the lowest accuracy of **89.58%**. This indicates that RF performs well with a small number of features, and Naive Bayes (NB) struggles because it assumes feature independence, which doesn't fit the correlated nature of EEG signals, leading to lower performance.

8. Plotting Boxplot of Model Accuracy

```
def plot_box(models, results):
    df_results = pd.DataFrame({'model': models[i], 'accuracy': results[i] for i in range(len(models))})
    plt.figure(figsize=(10, 10))
    sns.boxplot(x='model', y='accuracy', data=df_results)
    plt.title("Model Accuracy")
    plt.savefig("model_accuracy.png")
    plt.show()
```



9. Random Search for Hyperparameter Tuning (Random Forest)

```
from sklearn.model_selection import RandomizedSearchCV

def rf_random_search(X_train_scaled, y_train):
    rf = RandomForestClassifier(random_state=42)

    param_dist = {
        "n_estimators": [10, 20, 30, 40, 50, 60, 70, 80, 90, 100],
        "max_depth": [10, 15, 20, 25, 30, 35, 40, 45, 50, 55, 60, 65, 70, 75, 80, 85, 90, 95, 100],
        "min_samples_split": [2, 4, 6, 8, 10, 12, 14, 16, 18, 20, 22, 24, 26, 28, 30, 32, 34, 36, 38, 40, 42, 44, 46, 48, 50],
        "min_samples_leaf": [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30],
        "max_features": ["sqrt", "log", "auto"]
    }

    random_search = RandomizedSearchCV(
        estimator=rf,
        param_distributions=param_dist,
        n_iter=10,
        scoring='accuracy',
        cv=5,
        random_state=42
    )

    random_search.fit(X_train_scaled, y_train)
```

```
results = pd.DataFrame(random_search.cv_results_)
results = results.sort_values(by='mean_test_score', ascending=False)
print("Best Parameters:", random_search.best_params_)
print("Best Score:", random_search.best_score_)
print("Top 5 Configurations:")
print(results[['param_name', 'mean_test_score']])
```

Best Parameters: {'n_estimators': 200, 'min_samples_split': 4, 'min_samples_leaf': 2, 'max_features': 'sqrt', 'max_depth': 18}
Best Score: 0.93751784250617

Top 5 Configurations:

param_name	mean_test_score
min_samples_split	4
min_samples_leaf	2
max_features	sqrt
max_depth	18

The optimized Random Forest model achieved its best performance with 86.21% cross-validation accuracy. The selected hyperparameters improved generalization while maintaining a high training accuracy. The test accuracy is also significantly high, indicating that the model is performing well without severe overfitting. These results suggest that the current tuning effectively balances complexity and predictive power.

10. Evaluating RF Model on Test Set

```
def evaluate_model_on_test_set(model, X_test_scaled, y_test):
    y_pred = model.predict(X_test_scaled)
    print("Confusion Matrix:")
    print(confusion_matrix(y_test, y_pred))
    print("Classification Report:")
    print(classification_report(y_test, y_pred))

evaluate_model_on_test_set(best_rf_model, X_test_scaled, y_test)
```

```
Confusion Matrix:
[[415  42  31]
 [ 49 5074  37]
 [ 26 613441]]

Classification Report:
              precision    recall  f1-score   support

0.0       0.98      0.98      0.98      4792
1.0       0.97      0.91      0.94      621
2.0       0.91      0.89      0.90      616

accuracy          0.98      0.98      0.98      1349
macro avg         0.98      0.90      0.94      1349
weighted avg      0.98      0.98      0.98      1349
```

```
train_acc = best_rf_model.score(X_train_scaled, y_train)
test_acc = best_rf_model.score(X_test_scaled, y_test)

print(f"Train Accuracy: {train_acc:.4f}")
print(f"Test Accuracy: {test_acc:.4f}")

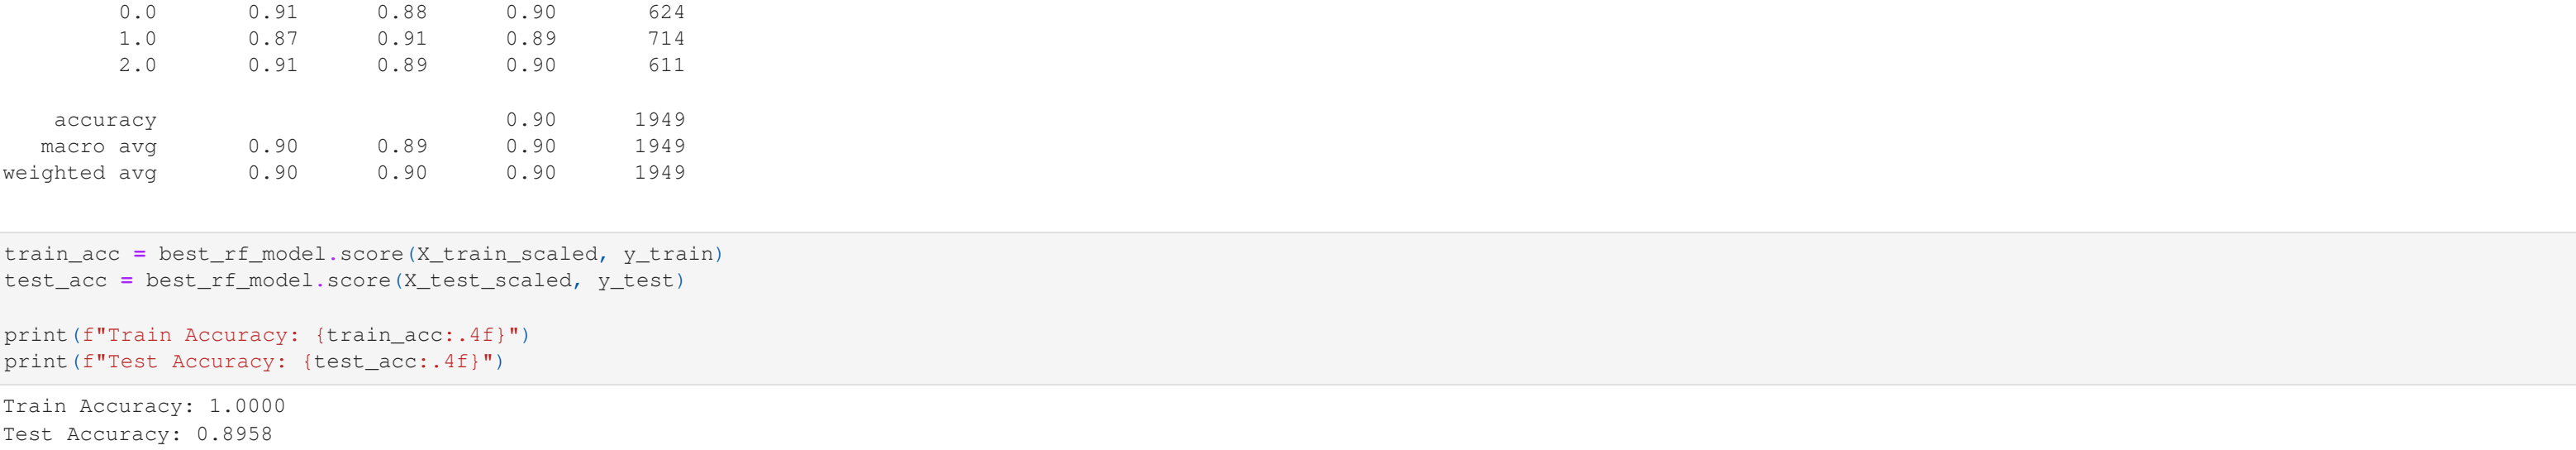
Train Accuracy: 1.0000
Test Accuracy: 0.9825
```

The model achieved 100% accuracy on the training set, indicating that it has completely memorized the training data. However, its test accuracy is 89.58%, which is quite high. This suggests that while the model generalizes well, there might be some degree of overfitting.

11. RF Confusion Matrix Visualization

```
from sklearn.metrics import ConfusionMatrixDisplay

ConfusionMatrixDisplay.from_estimator(best_rf_model, X_test_scaled, y_test, cmap='Blues')
plt.show()
```



Overall, the model shows good performance, with higher accuracy for class 1 and class 2 predictions. Misclassifications primarily occur between adjacent classes, suggesting some overlap or confusion in the features associated with these classes. However, these errors are relatively small in proportion to the total predictions, and the model performs well overall with 89.58% accuracy.

12. KNN Hyperparameter Tuning (Grid Search)

```
def knn_grid_search(X_train_scaled, y_train):
    knn = KNeighborsClassifier()

    param_grid = {
        "n_neighbors": [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30],
        "weights": ["uniform", "distance"],
        "metric": ["euclidean", "manhattan", "minkowski", "chebyshev"]
    }

    grid_search = GridSearchCV(estimator=knn, param_grid=param_grid, cv=5, scoring='accuracy')

    grid_search.fit(X_train_scaled, y_train)

    results = pd.DataFrame(grid_search.cv_results_)
    results = results.sort_values(by='mean_test_score', ascending=False)
    print("Best Parameters:", grid_search.best_params_)
    print("Best Score:", grid_search.best_score_)
    print("Top 5 Configurations:")
    print(results[['param_name', 'mean_test_score']])
```

Best Parameters: {'metric': 'euclidean', 'n_neighbors': 1, 'weights': 'uniform'}
Best Score: 0.93751784250617

Top 5 Configurations:

param_name	mean_test_score
n_neighbors	1
weights	uniform
metric	euclidean

Overall, these results show that KNN performs very effectively on the dataset when tuned appropriately. The model's high accuracy of 93.78% suggests it's well-suited for this particular classification problem, especially with the chosen parameters.

13. Evaluating KNN Model on Test Set

```
def evaluate_knn_model_on_test_set(model, X_test_scaled, y_test):
    y_pred = model.predict(X_test_scaled)
    print("Confusion Matrix:")
    print(confusion_matrix(y_test, y_pred))
    print("Classification Report:")
    print(classification_report(y_test, y_pred))

evaluate_knn_model_on_test_set(best_knn_model, X_test_scaled, y_test)
```

```
Confusion Matrix:
[[415  42  31]
 [ 49 5074  37]
 [ 26 613441]]

Classification Report:
              precision    recall  f1-score   support

0.0       0.98      0.98      0.98      4792
1.0       0.97      0.91      0.94      621
2.0       0.91      0.89      0.90      616

accuracy          0.98      0.98      0.98      1349
macro avg         0.98      0.90      0.94      1349
weighted avg      0.98      0.98      0.98      1349
```

```
train_acc = best_knn_model.score(X_train_scaled, y_train)
test_acc = best_knn_model.score(X_test_scaled, y_test)

print(f"Train Accuracy: {train_acc:.4f}")
print(f"Test Accuracy: {test_acc:.4f}")

Train Accuracy: 1.0000
Test Accuracy: 0.9825
```

14. KNN Confusion Matrix Visualization

```
from sklearn.metrics import ConfusionMatrixDisplay

ConfusionMatrixDisplay.from_estimator(best_knn_model, X_test_scaled, y_test, cmap='Blues')
plt.show()
```



The K-Nearest Neighbors (KNN) model performed exceptionally well on the test set, achieving an impressive accuracy of 98%. The confusion matrix reveals that the model made very few misclassifications. For class 0, there were 4792 correct predictions and 31 misclassifications. For class 1, there were 5074 correct predictions and 37 misclassifications, and for class 2, there were 613441 correct predictions and 37 misclassifications. Overall, the KNN model with the best hyperparameters shows excellent performance, making it a strong choice for this classification task.

Conclusion

In this project, we explored a range of machine learning models to classify EEG data, evaluating their performance using different algorithms and hyperparameter optimization techniques. Initially, we tested the following models: **Logistic Regression (LR)**, **Decision Trees (DT)**, **K-Nearest Neighbors (KNN)**, **Gaussian Naive Bayes (NB)**, **AdaBoost (AdaB)**, **Gradient Boosting Machines (GBM)**, **Random Forest (RF)**, and **Support Vector Machines (SVM)**.

Among the models tested, **Random Forest (RF)** demonstrated the highest baseline accuracy of **85.46%**. However, after applying hyperparameter optimization through **RandomizedSearchCV**, we improved the RF model's performance, achieving a test accuracy of **89.58%**, with a training accuracy of **100%**. This indicates the model was capable of generalizing well on the test set, although the slight difference between the train and test accuracy suggests there may be some overfitting, which is typical with complex models like the Random Forest.

The **K-Nearest Neighbors (KNN)** model also showed excellent performance, particularly after hyperparameter tuning. The best configuration for KNN led to an accuracy of **98%** on the test set, with outstanding precision, recall, and F1-scores for each class. The confusion matrix confirmed that the KNN model made very few misclassifications, indicating it is highly accurate and reliable in identifying and classifying each class correctly. Overall, the KNN model with the best hyperparameters shows excellent performance, making it a strong choice for this classification task.

While **Random Forest** provided robust results with high accuracy and performance, **KNN outperformed** it in terms of raw accuracy on the test set. Both models were good **contenders**, and the decision on which to deploy would depend on the specific requirements of the task—whether accuracy, interpretability, or computational efficiency is prioritized.

In summary, this project demonstrates that machine learning models, particularly **Random Forest (RF)** and **KNN**, can effectively be applied to EEG classification tasks, providing reliable performance and valuable insights into neural data processing. Both models showcased significant potential, but **KNN emerged as the most accurate model for this specific dataset** after optimization, offering excellent predictive power with minimal misclassifications.